

Module 3 Assignment | Analysis with MapReduce/Pig/Hive

Trang Tran

CPS, Northeastern University

ALY6110 | Data Management and Big Data

Professor Andrew Kinley

Aug 08, 2023

Set up the directory and subfolders in HDFS.

```
hadoop fs -mkdir lab
```

```
hadoop fs -mkdir lab/input
```

```
hadoop fs -mkdir lab/output
```

```
#transfer files into HDFS
```

```
hadoop fs -copyFromLocal orders.csv lab/input
```

```
hadoop fs -copyFromLocal inventory.csv lab/input
```

```
nano inventory.csv #then delete the header and save it back into the file
```

```
cat inventory.csv
```

Pig # launches the Pig shell

```
1. orders = LOAD 'lab/input/orders.csv' USING PigStorage(',') AS (table_id: chararray, order_id: chararray, product_id: int, date: chararray);
```

Load the 'orders.csv' file using PigStorage with comma as the delimiter, also assign aliases to the columns, specifying the data types for each column: table_id as a chararray, order_id as a chararray, product_id as an integer, and date as a chararray.

```
2. inventory = LOAD 'lab/input/inventory.csv' USING PigStorage(',') AS (table_id: chararray, product_id: int, name: chararray, category: chararray, inventory: int);
```

Load the 'inventory.csv' file using PigStorage with comma as the delimiter, also assign aliases to the columns, specifying the data types for each column: table_id as a chararray, product_id as an integer, name as a chararray, category as a chararray, and inventory as an integer.

```
3. grouped_orders = GROUP orders BY product_id;
```

Group the 'orders' table by the 'product_id' column, and store the result in the new table named 'grouped_orders'.

4. `orders_count = FOREACH grouped_orders GENERATE group AS product_id,
COUNT(orders) AS order_count;`

For each group in 'grouped_orders', calculate the count of orders for each 'product_id', then generate a relation named 'orders_count' with columns product_id and order_count.

5. `grouped_inventory = GROUP inventory BY product_id;`

Group the 'inventory' relation by the 'product_id' column, and store the result in the new table named 'grouped_inventory'.

6. `inventory_sum = FOREACH grouped_inventory GENERATE group AS product_id,
SUM(inventory.inventory) AS total_inventory;`

For each group in 'grouped_inventory', calculate the sum of 'inventory' values for each product_id, then generate a relation 'inventory_sum' with columns product_id and total_inventory.

7. `joined_data = JOIN orders_count BY product_id LEFT OUTER, inventory_sum BY
product_id;`

Perform a left join between orders_count and inventory_sum based on the 'product_id' column, then store a relation 'joined_data' with columns from both orders_count and inventory_sum.

8. `joined_data = FOREACH joined_data GENERATE orders_count::product_id AS product_id,
orders_count::order_count AS order_count, inventory_sum::total_inventory AS total_inventory,
(order_count - total_inventory) AS difference;`

For each row in 'joined_data', calculate the difference between the order_count and total_inventory columns, and generate a new column named 'difference'.

9. `insufficient_stock = FILTER joined_data BY difference > 0;`

Filter the rows in 'joined_data' to retain only those where the difference column is greater than 0 (where order_count > total_inventory), indicating insufficient stock for orders.

10. result = FOREACH insufficient_stock GENERATE product_id, difference AS additional_items;

For each row in 'insufficient_stock', generate the product_id and additional_items columns, where additional_items is the value of the difference column, then store it in the 'result' table (shown below).

HIVE

```
hadoop fs -mkdir HV
```

```
hadoop fs -mkdir HV/input
```

#transfer files into the directory

```
hadoop fs -copyFromLocal orders.csv HV/input
```

```
hadoop fs -copyFromLocal inventory.csv HV/input
```

Hive

1. CREATE TABLE orders (table_id STRING, order_id STRING, product_id INT) ROW FORMAT DELIMITED FIELDS TERMINATED BY ',' STORED AS TEXTFILE;

Create a new blank table named 'orders' with three columns: table_id of type STRING, order_id of type STRING, and product_id of type INT. The data in this table will be delimited by commas, and it will be stored in text file format.

2. LOAD DATA INPATH 'hdfs:///user/hadoop/HV/input/orders.csv' OVERWRITE INTO TABLE orders;

Load data from the specified file 'orders.csv' from the HDFS HV/input directory into the 'orders' table.

then we do the same thing with the 'inventory' table in lines 3 and 4 below

```
3. CREATE TABLE inventory (table_id STRING, product_id INT, name STRING, category
STRING, inventory INT) ROW FORMAT DELIMITED FIELDS TERMINATED BY ','
STORED AS TEXTFILE;
```

```
4. LOAD DATA INPATH 'hdfs:///user/hadoop/HV/input/inventory.csv' OVERWRITE INTO
TABLE inventory;
```

```
5. CREATE TABLE orders_count AS SELECT product_id, COUNT(order_id) AS order_count
FROM orders GROUP BY product_id;
```

Create a new table named 'orders_count' by selecting the 'product_id' column and calculate the count of orders for each product_id and assign the result to the 'order_count' column.

```
6. CREATE TABLE inventory_sum AS SELECT product_id, SUM(inventory) AS
total_inventory FROM inventory GROUP BY product_id;
```

Create a new table named 'inventory_sum' by calculating the sum of inventory values for each 'product_id' and assign the result to the 'total_inventory' column.

```
7. CREATE TABLE joined_data AS SELECT o.product_id, o.order_count, i.total_inventory,
(o.order_count - i.total_inventory) AS difference FROM orders_count o LEFT JOIN
inventory_sum i ON o.product_id = i.product_id;
```

Create a new table named 'joined_data' by performing a LEFT JOIN between the 'orders_count' and 'inventory_sum' tables by the 'product_id' column.

```
8. CREATE TABLE insufficient_stock AS SELECT product_id, difference AS additional_items
FROM joined_data WHERE difference > 0;
```

Create a new table named 'insufficient_stock' by filtering and selecting rows where the difference column is greater than 0 and rename the difference column as 'additional_items'.

9. SELECT * FROM insufficient_stock;

View all content in the 'insufficient_stock' table (shown below).

(35240,1)	(187117,1)	hive> select * from insufficient_stock;
(35242,1)	(202497,1)	OK
(35270,1)	(228114,1)	35274 1
(35274,1)	(231571,1)	35362 8
(35291,2)	(232959,116)	35387 1
(35308,1)	(239984,1)	35393 1
(35362,8)	(244060,1)	35516 158
(35387,1)	(282384,1)	35538 26
(35393,1)	(306207,1)	35579 27
(35405,1)	(312623,6)	35585 2
(35516,158)	(312624,623)	35915 664
(35518,83)	(312626,7)	36444 1
(35519,1)	(312634,4)	37195 1
(35537,5)	(312638,1)	38014 1
(35538,26)	(312651,1)	43501 1
(35554,675)	(312652,1)	46897 1
(35564,2)	(312661,9)	47626 1
(35579,27)	(312663,31)	50613 1
(35585,2)	(312666,1)	53721 1
(35915,664)	(312675,7)	54183 1
(36422,1)	(312681,1)	57668 1
(36444,1)	(312741,8)	58930 2
(36912,1)	(312746,3)	61731 1
(37149,1)	(312752,1)	65026 1
(37195,1)	(312760,1)	71334 1
(37262,1)	(312775,1)	72611 1
(38014,1)	(312788,1)	75080 1
(38018,1)	(312808,10)	75091 1
(42072,1)	(312812,11)	81260 1
(43501,1)	(312817,1)	86954 1
(45919,1)	(312937,1)	202497 1
(46897,1)	(312960,105)	228114 1
(47127,1)	(313049,1072)	239984 1
(47626,1)	(313052,6)	282384 1
(50613,1)	(313078,1)	306207 1
(52220,1)	(313083,4)	312638 1
(52221,1)	(313084,1)	312651 1
(53721,1)	(313120,1)	312652 1
(54183,1)	(313135,180)	312661 9
(55185,1)	(313136,1)	312666 1
(57102,1)	(313143,1803)	312675 7
(57668,1)	(313151,1)	312752 1
(58930,2)	(389401,2)	312760 1
(60792,1)	(389568,1)	312775 1
(61731,1)	(389589,10)	312808 10
(65026,1)	(389624,399)	312812 11
(65028,1)	(389626,4)	312960 105
(67360,1)	(389643,1)	313049 1072
(71271,1)	(389942,2)	313136 1
(71334,1)	(390025,38)	313143 1803
(72611,1)	(390603,4)	313151 1
(74038,1)	(392382,10)	389589 10
(74470,1)	(392633,73)	389626 4
(75080,1)	(393019,1)	389643 1
(75091,1)	(393020,18)	389942 2
(75424,1)	(393085,6)	390025 38
(76081,1)	(393286,1)	390603 4
(76633,1)	(393309,1254)	392633 73
(78706,1)	(393738,24)	393019 1
(81260,1)	(393739,259)	393020 18
(83133,1)	(400454,8)	393085 6
(86954,1)	(400458,6)	393738 24
(88960,1)	(400468,15)	393739 259
(95749,1)	(400477,1)	400454 8
	(400479,2)	400479 2
	(402184,1)	402184 1
	(402189,1)	402189 1
	(402190,2)	402190 2
	(402195,1)	402198 21
	(402198,21)	402209 2
	(402209,2)	402475 1
	(402212,6)	402492 1
	(402243,1)	35240 1
	(402475,1)	35242 1
	(402492,1)	35270 1

MapReduce

```
hadoop fs -mkdir lab/MR
```

```
hadoop fs -copyFromLocal combined.csv lab/MR
```

```
hadoop fs -copyFromLocal inventory_mapper.py lab/MR
```

```
hadoop fs -copyFromLocal inventory_reducer.py lab/MR
```

```
chmod +x inventory_mapper.py
```

```
chmod +x inventory_reducer.py
```

```
hadoop jar /usr/lib/hadoop/hadoop-streaming.jar -files
```

```
inventory_mapper.py,inventory_reducer.py -mapper inventory_mapper.py -reducer
```

```
inventory_reducer.py -input lab/MR/combined.csv -output lab/MR/output
```

Mapper

```
# import the module for reading and writing data
```

```
import sys
```

```
# input is read by STDIN (standard input) and do the following for each input line
```

```
for line in sys.stdin:
```

```
    # remove leading and trailing whitespace
```

```
    line = line.strip()
```

```
    # split the line by comma separator, a list is produced
```

```
    line = line.split(",")
```

```
    # assign values of the list to the variable names
```

```
    table_name = line[0]
```

```
    product_id = line[1]
```

```
    inventory_value = line[-1] # Assuming inventory value is the last column
```

```

if table_name == 'order_id':

    print('%s\t%s' % (product_id, 'order'))

elif table_name == 'inventory_id':

    print('%s\t%s' % (product_id, inventory_value))

```

Reducer

```

# import the module for reading and writing data

import sys

# Initialize variables: lastKey as None and maxValue as 0

(lastKey, maxValue) = (None, 0)

# Read input from STDIN (standard input) line by line

for line in sys.stdin:

    # Remove leading and trailing whitespace

    line = line.strip()

    # Split the line into key and value based on tab separation

    (key, value) = line.split('\t', 1)

    # Convert the value to an integer

    value = int(value)

    # Check if lastKey is defined and not equal to the current key

    if lastKey and lastKey != key:

        # Output the product_id and the difference between inventory and orders

        print('%s\t%s' % (lastKey, maxValue - value))

        (lastKey, maxValue) = (key, value)

```



```
else:
```

```
# Update maxValue if the current value is larger
```

```
maxValue = max(maxValue, value)
```

```
lastKey = key
```

```
# Output the result for the last product_id
```

```
if lastKey:
```

```
print('%s\t%s' % (lastKey, maxValue))
```

```
[[hadoop@ip-172-31-60-70 ~]$ hadoop jar /usr/lib/hadoop/hadoop-streaming.jar -files inventory_mapper.py,inventory_reducer.py -mapper inventory_mapper.py -reducer inventory_reducer.py -input lab/MR/combined.csv -output lab/MR/output6
packageJobJar: [ ] [/usr/lib/hadoop/hadoop-streaming-3.3.3-amzn-4.jar] /tmp/streamjob7107236983720592866.jar tmpDir=null
2023-08-11 03:38:26,863 INFO client.DefaultNoHARMAFailoverProxyProvider: Connecting to ResourceManager at ip-172-31-60-70.ec2.internal/172.31.60.70:8032
2023-08-11 03:38:27,118 INFO client.AHSProxy: Connecting to Application History server at ip-172-31-60-70.ec2.internal/172.31.60.70:10200
2023-08-11 03:38:27,177 INFO client.DefaultNoHARMAFailoverProxyProvider: Connecting to ResourceManager at ip-172-31-60-70.ec2.internal/172.31.60.70:8032
2023-08-11 03:38:27,178 INFO client.AHSProxy: Connecting to Application History server at ip-172-31-60-70.ec2.internal/172.31.60.70:10200
2023-08-11 03:38:27,499 INFO mapreduce.JobResourceUploader: Disabling Erasure Coding for path: /tmp/hadoop-yarn/staging/hadoop/.staging/job_1691721480693_0003
2023-08-11 03:38:27,899 INFO lzo.GPLNativeCodeLoader: Loaded native gpl library
2023-08-11 03:38:27,903 INFO lzo.LzoCodec: Successfully loaded & initialized native-lzo library [hadoop-lzo rev 049362b7cf53ff5f739d6b1532457f2c6cd495e8]
2023-08-11 03:38:27,945 INFO mapred.FileInputFormat: Total input files to process : 1
2023-08-11 03:38:28,003 INFO mapreduce.JobSubmitter: number of splits:4
2023-08-11 03:38:28,258 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1691721480693_0003
2023-08-11 03:38:28,258 INFO mapreduce.JobSubmitter: Executing with tokens: [ ]
2023-08-11 03:38:28,592 INFO conf.Configuration: resource-types.xml not found
2023-08-11 03:38:28,593 INFO resource.ResourceUtils: Unable to find 'resource-types.xml'.
2023-08-11 03:38:28,756 INFO impl.YarnClientImpl: Submitted application application_1691721480693_0003
2023-08-11 03:38:28,848 INFO mapreduce.Job: The url to track the job: http://ip-172-31-60-70.ec2.internal:20888/proxy/application_1691721480693_0003/
2023-08-11 03:38:28,851 INFO mapreduce.Job: Running job: job_1691721480693_0003
2023-08-11 03:38:38,014 INFO mapreduce.Job: Job job_1691721480693_0003 running in uber mode : false
2023-08-11 03:38:38,015 INFO mapreduce.Job: map 0% reduce 0%
2023-08-11 03:38:51,177 INFO mapreduce.Job: map 50% reduce 0%
2023-08-11 03:39:03,280 INFO mapreduce.Job: map 100% reduce 0%
2023-08-11 03:39:07,306 INFO mapreduce.Job: Task Id : attempt_1691721480693_0003_r_000000_0, Status : FAILED
Error: java.lang.RuntimeException: PipeMapRed.waitOutputThreads(): subprocess failed with code 1
    at org.apache.hadoop.streaming.PipeMapRed.waitOutputThreads(PipeMapRed.java:326)
    at org.apache.hadoop.streaming.PipeMapRed.mapRedFinished(PipeMapRed.java:539)
    at org.apache.hadoop.streaming.PipeReducer.reduce(PipeReducer.java:128)
    at org.apache.hadoop.mapred.ReduceTask.runOldReducer(ReduceTask.java:446)
    at org.apache.hadoop.mapred.ReduceTask.run(ReduceTask.java:394)
    at org.apache.hadoop.mapred.YarnChild$2.run(YarnChild.java:178)
    at java.security.AccessController.doPrivileged(Native Method)
    at javax.security.auth.Subject.doAs(Subject.java:422)
    at org.apache.hadoop.security.UserGroupInformation.doAs(UserGroupInformation.java:1878)
    at org.apache.hadoop.mapred.YarnChild.main(YarnChild.java:172)
2023-08-11 03:39:14,463 INFO mapreduce.Job: Task Id : attempt_1691721480693_0003_r_000000_1, Status : FAILED
Error: java.lang.RuntimeException: PipeMapRed.waitOutputThreads(): subprocess failed with code 1
    at org.apache.hadoop.streaming.PipeMapRed.waitOutputThreads(PipeMapRed.java:326)
    at org.apache.hadoop.streaming.PipeMapRed.mapRedFinished(PipeMapRed.java:539)
    at org.apache.hadoop.streaming.PipeReducer.reduce(PipeReducer.java:128)
    at org.apache.hadoop.mapred.ReduceTask.runOldReducer(ReduceTask.java:446)
    at org.apache.hadoop.mapred.ReduceTask.run(ReduceTask.java:394)
    at org.apache.hadoop.mapred.YarnChild$2.run(YarnChild.java:178)
    at java.security.AccessController.doPrivileged(Native Method)
    at javax.security.auth.Subject.doAs(Subject.java:422)
    at org.apache.hadoop.security.UserGroupInformation.doAs(UserGroupInformation.java:1878)
    at org.apache.hadoop.mapred.YarnChild.main(YarnChild.java:172)
2023-08-11 03:39:22,513 INFO mapreduce.Job: Task Id : attempt_1691721480693_0003_r_000000_2, Status : FAILED
Error: java.lang.RuntimeException: PipeMapRed.waitOutputThreads(): subprocess failed with code 1
    at org.apache.hadoop.streaming.PipeMapRed.waitOutputThreads(PipeMapRed.java:326)
    at org.apache.hadoop.streaming.PipeMapRed.mapRedFinished(PipeMapRed.java:539)
    at org.apache.hadoop.streaming.PipeReducer.reduce(PipeReducer.java:128)
    at org.apache.hadoop.mapred.ReduceTask.runOldReducer(ReduceTask.java:446)
    at org.apache.hadoop.mapred.ReduceTask.run(ReduceTask.java:394)
    at org.apache.hadoop.mapred.YarnChild$2.run(YarnChild.java:178)
    at java.security.AccessController.doPrivileged(Native Method)
    at javax.security.auth.Subject.doAs(Subject.java:422)
    at org.apache.hadoop.security.UserGroupInformation.doAs(UserGroupInformation.java:1878)
    at org.apache.hadoop.mapred.YarnChild.main(YarnChild.java:172)
```

****I followed all the steps above and used 2 Python files as stated for 'Mapper' and 'Reducer', but I still couldn't go through the final streaming result. I tried editing the script several times and run it again and again but still cannot figure out the errors in this case.*